EXPLORING FINE-TUNING STRATEGIES FOR DISTILBERT GENERALIZATION

Ziyi Liu

HELLO, MEET YOUR HOST!

DISTIL BERT

PLAYER 01

SMALL DISTILBERT

PLAYER 02

SMALL DISTILBERT + LORA

PLAYER 03

HANDSOME YOUNG MAN _ DISTILBERT + LORA

HOW TO PLAY

01. Train, validate, test on SST-5

02. Fixed epochs, learning rate, batch size.

03. Best accuracy and confusion matrix win.

04. Faster model gets +10 points.

05. No cheating on computational resources.

STANFORD SENTIMENT TREEBANK

01. 11,855 single sentences extracted from movie reviews.

02. With 5 labels: very positive, positive, neutral, negative, very negative

03. Train: 8.54 K rows

Val: 1.1 K rows

Test: 2.21 K rows

LET'S PLAY

HOST – DISTIL BERT

'TEST_ACCURACY': 50.31%

PLAYER01– SMALL DISTIL BERT

'TEST_ACCURACY': 39.82%

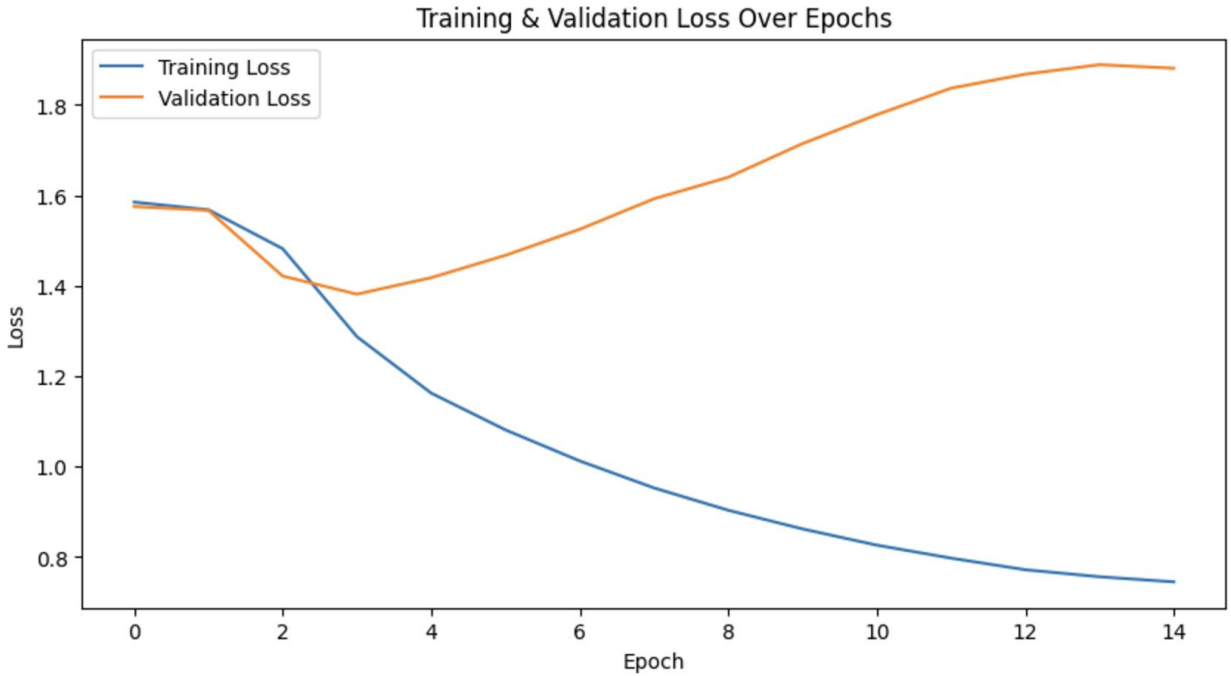

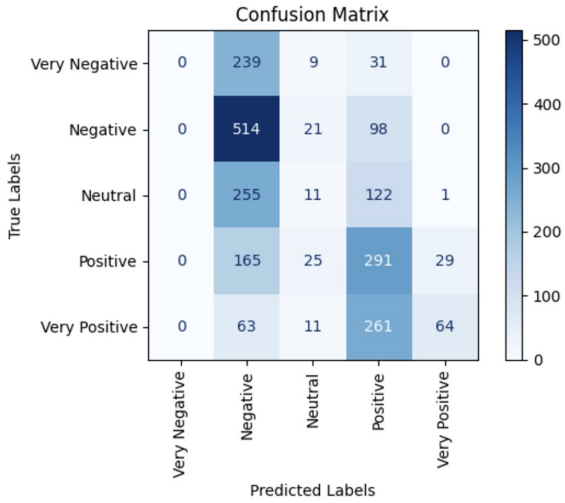

PLAYER02– SMALL DISTIL BERT + LORA

'TEST_ACCURACY': 28.05%

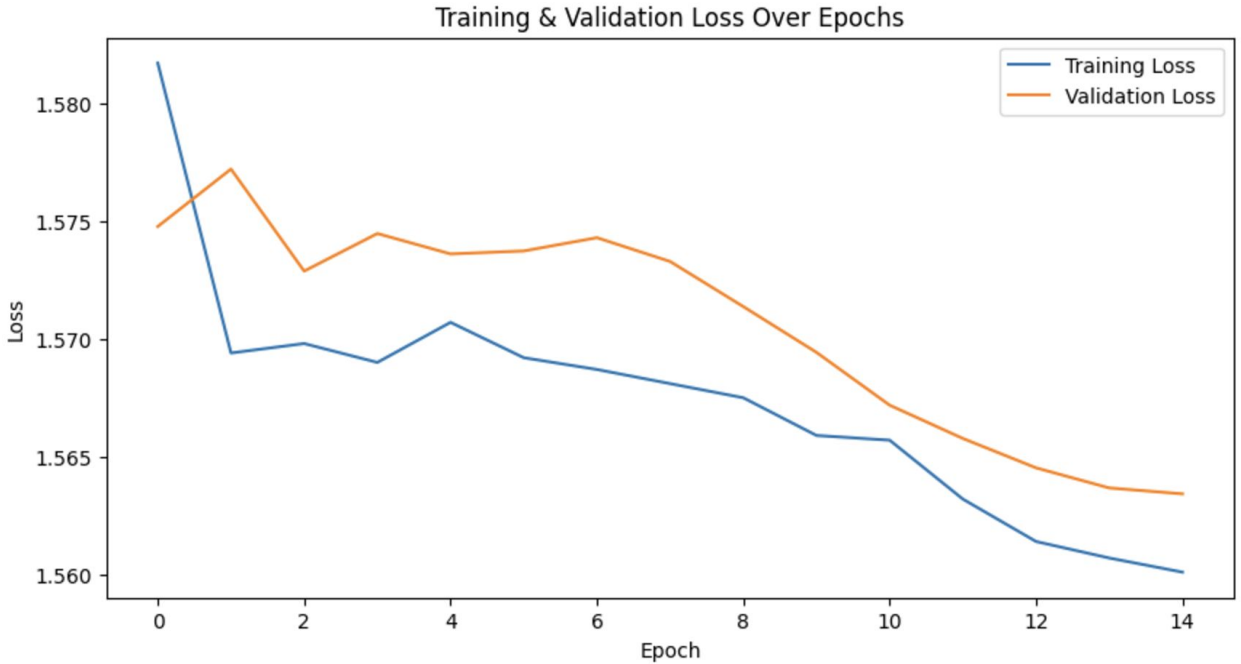

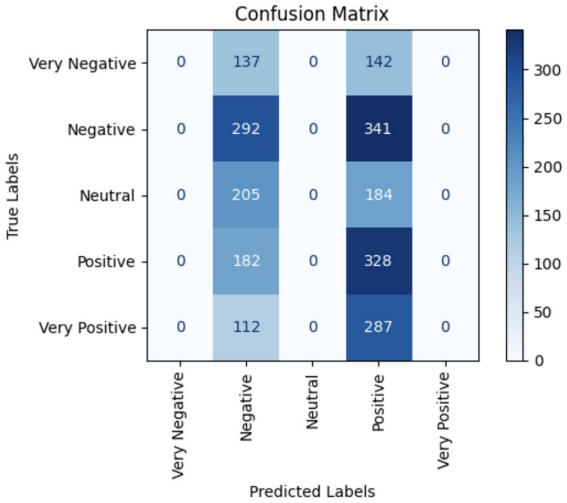

PLAYER03– DISTIL BERT + LORA

'TEST_ACCURACY': 50.54%

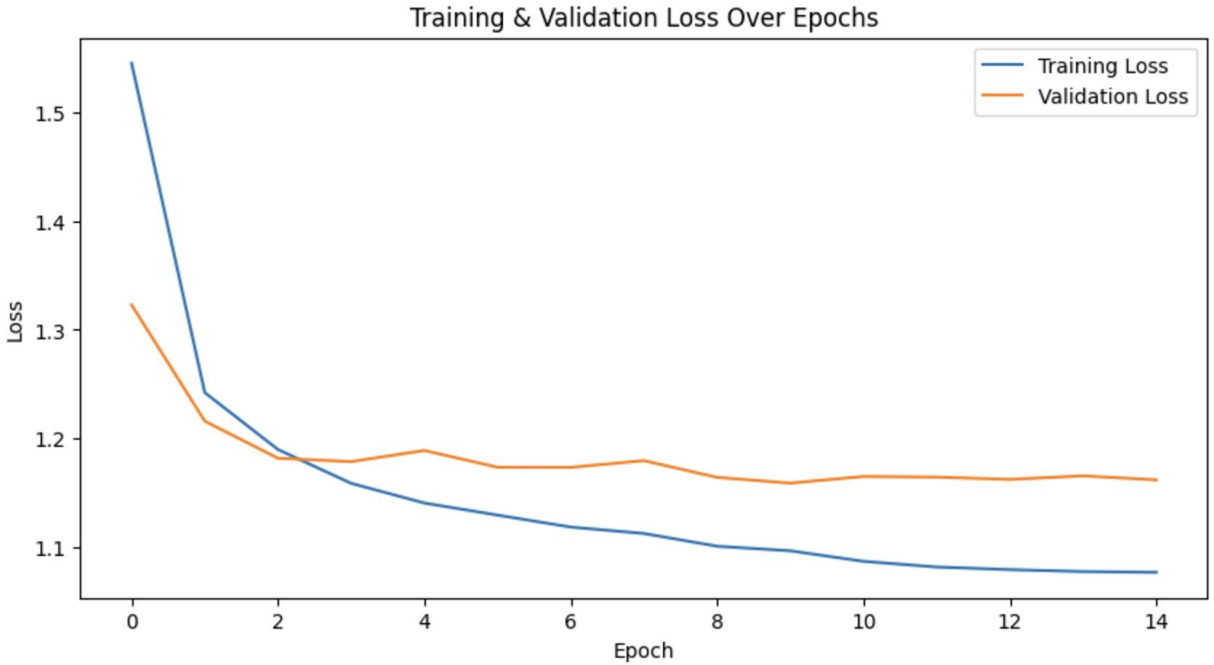

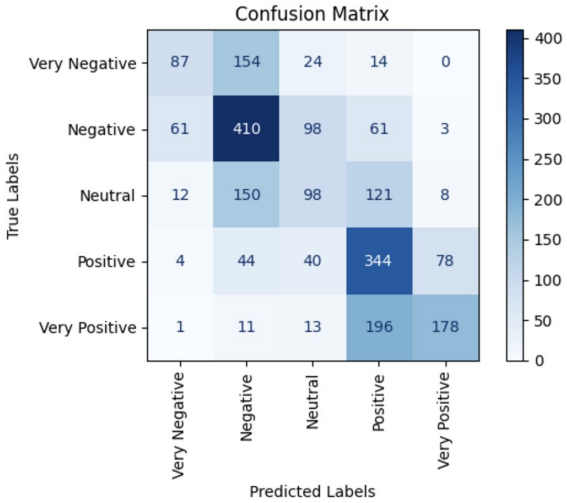

SUBMIT YOUR VOTE


HANDSOME YOUNG MAN _
DISTILBERT + LORA

STICKER SHEET